



SaltStack as network orchestrator

Scalable, fast, cross-vendor

Mircea Ulinic
Cloudflare, London

APRICOT 2017
Ho Chi Minh City, VN

Agenda

- So you want to automate
- Meet the tools
- Configure SaltStack
- CLI syntax
- Configuration management (brief intro)
- Real-world orchestration example

To automate, I have to learn Python or another programming language.


To automate, I have to learn Python or another programming language.

WRONG!

Do not jump into implementation.
Design first!

What's the best tool?

Wrong question.



What's the best tool?

What's the best tool for my network?

What's the best tool for my network?

- Mind your network
- How many devices?
- How many platforms / operating systems?
- How dynamic?
- Configuration management only?
- Triggered configuration changes?
- External sources of truth? e.g. IPAM
- Do you need native caching? REST API?
etc...

Meet the Tools

Why Salt?

- Very scalable
- Concurrency
- Easily configurable & customizable
- Config verification & enforcement
- Periodically collect statistics
- Native caching and drivers for useful tools

Meet the Tools

Why Salt?

“

In SaltStack, speed isn't a byproduct, it is a design goal. SaltStack was created as an extremely fast, lightweight communication bus to provide the foundation for a remote execution engine.

SaltStack now provides orchestration, configuration management, event reactors, cloud provisioning, and more, all built around the SaltStack high-speed communication bus.


... + cross-vendor network automation from 2016.11 (Carbon)

”

<https://docs.saltstack.com/en/getstarted/speed.html>

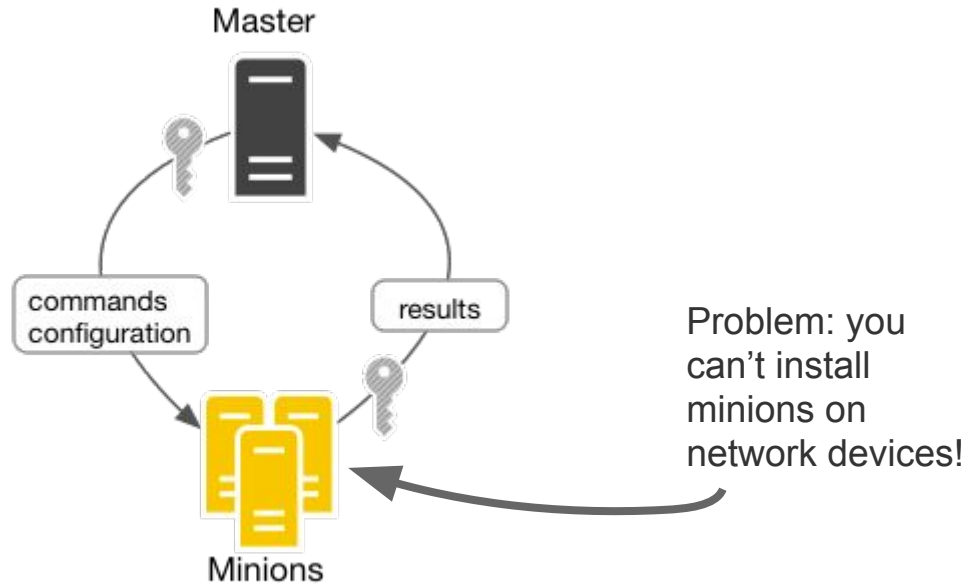
Meet the Tools

Orchestration vs. Automation



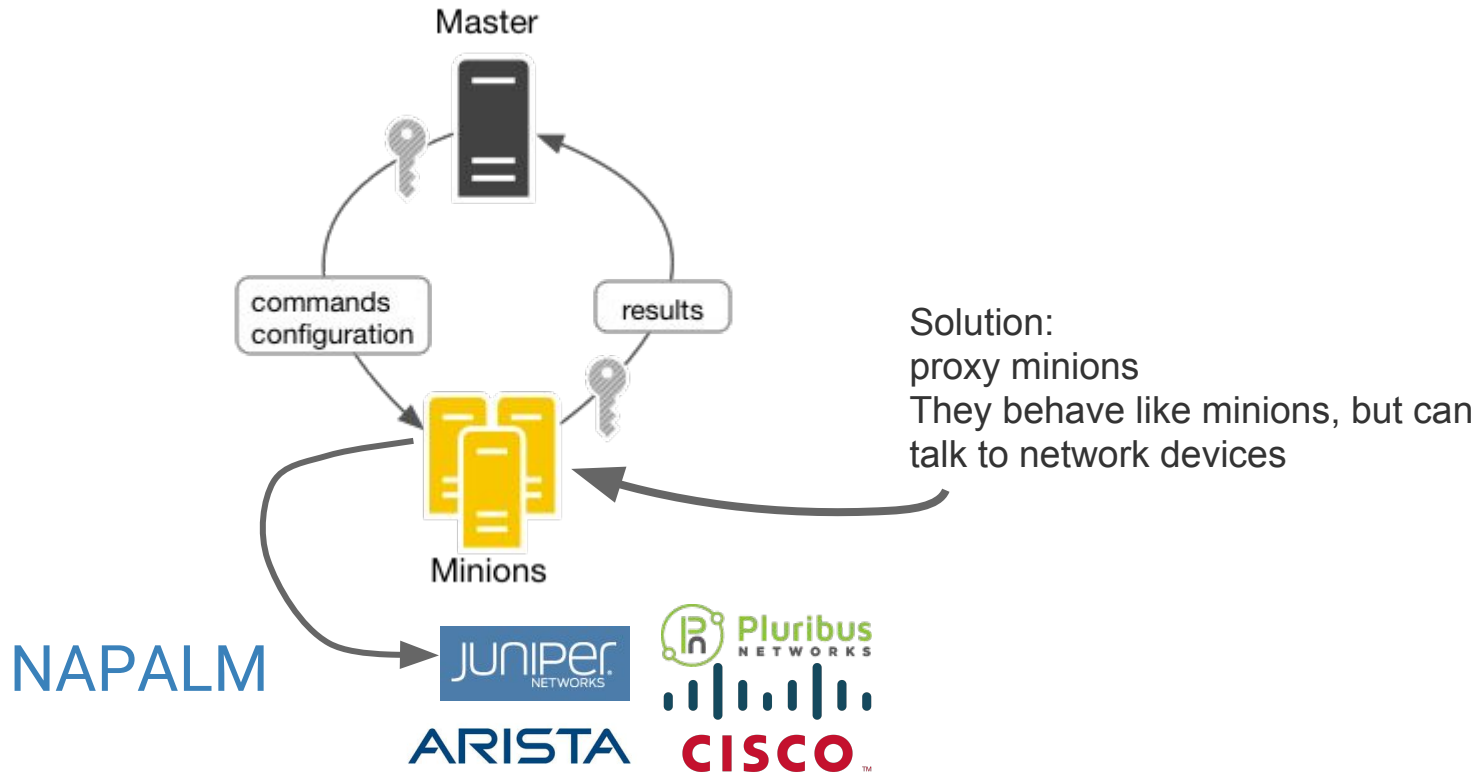
Meet the Tools

Salt Architecture



Meet the Tools


Salt Architecture



Meet the Tools

Why NAPALM?

(Network Automation and Programmability Abstraction Layer with Multivendor support)



<https://github.com/napalm-automation>



NAPALM integrated in SaltStack

NETWORK AUTOMATION: NAPALM

Beginning with 2016.11.0, network automation is included by default in the core of Salt. It is based on the [NAPALM](#) library and provides facilities to manage the configuration and retrieve data from network devices running widely used operating systems such as: JunOS, IOS-XR, eOS, IOS, NX-OS etc. - see [the complete list of supported devices](#).

The connection is established via the `NAPALM proxy`.

In the current release, the following modules were included:

- `NAPALM grains` - Select network devices based on their characteristics
- `NET execution module` - Networking basic features
- `NTP execution module`
- `BGP execution module`
- `Routes execution module`
- `SNMP execution module`
- `Users execution module`
- `Probes execution module`
- `NTP peers management state`
- `SNMP configuration management state`
- `Users management state`

<https://docs.saltstack.com/en/develop/topics/releases/2016.11.0.html>

Configure SaltStack

New to Salt?

Pillar

Free-form data that can be used to organize configuration values or manage sensitive data, e.g.: interface details, NTP peers, BGP config...

YAML file / database / git repository ... etc.

Grains

data collected from the device, e.g.: device model, vendor, uptime, serial number etc.

Salt handles this, you don't need to do anything

Salt in 10 minutes: <https://docs.saltstack.com/en/latest/topics/tutorials/walkthrough.html>

Configure SaltStack

Master config

/etc/salt/master

```
file_roots:  
  base:  
    - /etc/salt/states  
    - /etc/salt/reactors  
    - /etc/salt/templates  
pillar_roots:  
  base:  
    - /etc/salt/pillar
```

Environment name

Useful to have different environments: prod, qa, develop etc.

For the beginning, let's focus only on **file_roots** and **pillar_roots**. The others settings are more advanced features: <https://docs.saltstack.com/en/latest/ref/configuration/master.html>

[Complete salt master config file](#)

Configure SaltStack

Device *pillar*

Under the **pillar_roots** directory (as configured in `/etc/salt/master`):

`/etc/salt/pillar/device1.sls`

```
proxy:
  proxytype: napalm
  driver: junos
  host: hostname_or_ip_address
  username: my_username
  passwd: my_password
```

Mandatory

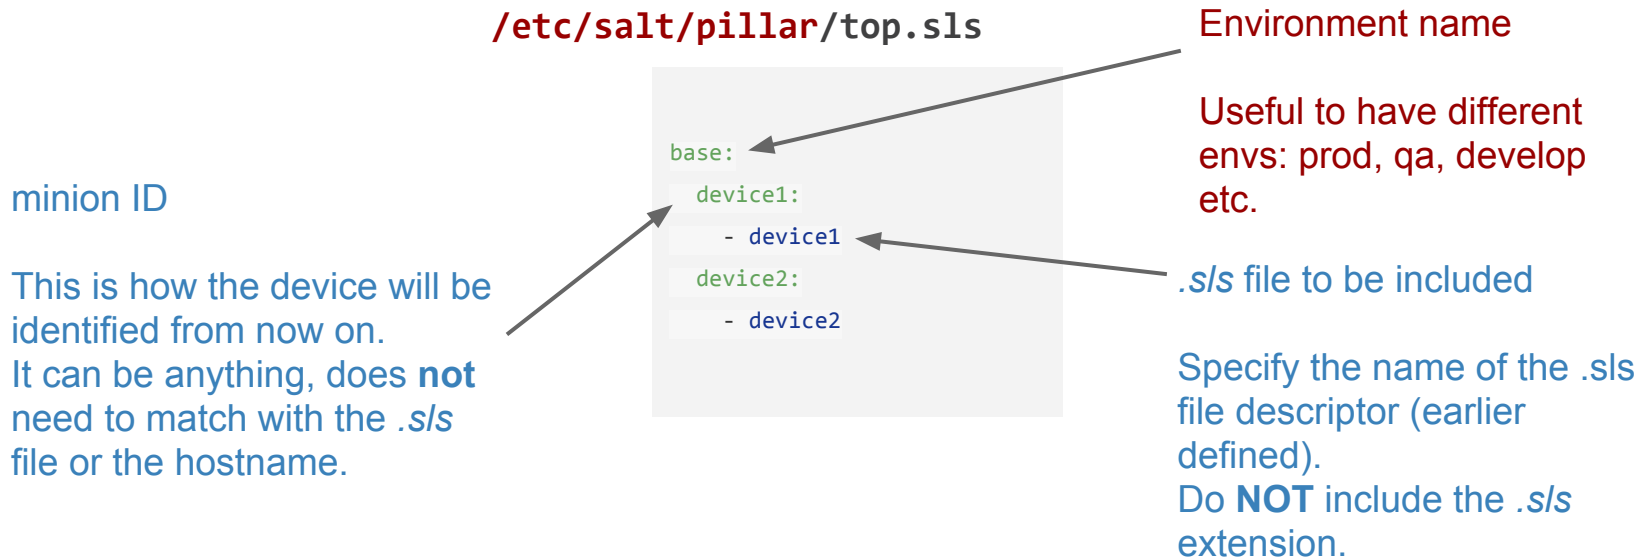
Choose between: junos,
eos, ios, iosxr, nxos, etc.
See the **complete list**.

Complete documentation at: <https://docs.saltstack.com/en/develop/ref/proxy/all/salt.proxy.napalm.html>

Configure SaltStack

The *top* file

Under the **pillar_roots** directory (as configured in `/etc/salt/master`):



Salt CLI syntax

Selecting the devices we need to run the command.

Targeting can be complex:

<https://docs.saltstack.com/en/latest/topics/targeting/>



```
$ sudo salt <target> <function> [<arguments>]
```

The diagram illustrates the Salt CLI syntax with arrows pointing from explanatory text to the components of the command. An arrow points from the URL 'https://docs.saltstack.com/en/latest/topics/targeting/' to the '<target>' placeholder. Another arrow points from the text 'Function name, as specified in the module documentation.' to the '<function>' placeholder. A third arrow points from the text 'Function arguments, as specified in the module documentation. Some functions do not require any arguments.' to the '[<arguments>]' placeholder.

Function name, as specified in the module documentation.

For example if we need BGP-related commands, we'll look at the [BGP module](#).

Other examples: [dnsutil.A](#), [net.arp](#), [net.ldap](#), [net.traceroute](#) etc.

Function arguments, as specified in the module documentation.
Some functions do not require any arguments.

Salt CLI syntax

Examples

```
$ sudo salt 'edge*' net.traceroute 8.8.8.8
# execute traceroute on all devices whose minion ID starts with 'edge'

$ sudo salt -N NA transit.disable cogent
# disable Cogent in North-America

$ sudo salt -G 'os:junos' net.cli "show version"
# execute 'show version' on all devices running JunOS

$ sudo salt -C 'edge* and G@os:iosxr and G@version:6.0.2' net.arp
# get the ARP tables from devices whose ID starts with edge*, running IOS-XR 6.0.2

$ sudo salt -G 'model:MX480' probes.results
# retrieve the results of the RPM probes from all Juniper MX480 routers
```

→ 'NA' is a nodegroup:

<https://docs.saltstack.com/en/latest/topics/targeting/nodegroups.html>

Salt CLI syntax

Output example

Default output style: nested.

```
$ sudo salt edge01.iad01 net.arp
```

```
edge01.iad01:
```

```
out:
```

```
|_
```

```
age:
```

```
129.0
```

```
interface:
```

```
ae2.100
```

```
ip:
```

```
10.0.0.1
```

```
mac:
```

```
00:0f:53:36:e4:50
```

```
|_
```

```
age:
```

```
1101.0
```

```
interface:
```

```
xe-0/0/3.0
```

```
ip:
```

```
10.0.0.2
```

```
mac:
```

```
00:1d:70:83:40:c0
```


Salt CLI syntax

Outputters

```
$ salt --out=json edge01.iad01 net.arp
```

```
[
  {
    "interface": "ae2.100",
    "ip": "10.0.0.1",
    "mac": "00:0f:53:36:e4:50",
    "age": 129.0
  },
  {
    "interface": "xe-0/0/3.0",
    "ip": "10.0.0.2",
    "mac": "00:1d:70:83:40:c0",
    "age": 1101.0
  },
]
```

Using the **--out** optional argument, one can select the output format.

```
$ salt --out=yaml edge01.iad01 net.arp
```

```
edge01.iad01:
  comment: ''
  out:
    - age: 129.0
      interface: ae2.100
      ip: 10.0.0.1
      mac: 00:0f:53:36:e4:50
    - age: 1101.0
      interface: xe-0/0/3.0
      ip: 10.0.0.2
      mac: 00:1d:70:83:40:c0
```

Configuration management

Cross vendor templating (1)

/etc/salt/templates/example.jinja

Hostname taken
from the pillar.

```
{%- set router_vendor = grains.vendor -%}  
{%- set hostname = pillar.proxy.host -%}  
{%- if router_vendor|lower == 'juniper' %}  
system {  
    host-name {{ hostname }}.lab;  
}  
{%- elif router_vendor|lower in ['cisco', 'arista'] %}  
{# both Cisco and Arista have the same syntax for hostname #}  
hostname {{ hostname }}.lab  
{%- endif %}
```

Get the device
vendor from the
grains.

Multiple templating systems supported (not only Jinja):

<https://docs.saltstack.com/en/latest/ref/renderers/all/index.html>

Configuration management

Cross vendor templating (2)

```
$ sudo salt '*' net.load_template /etc/salt/templates/example.jinja
```

```
edge01.bjm01:
```

```
-----
```

```
already_configured:
```

```
False
```

```
comment:
```

```
diff:
```

```
@@ -35,7 +35,7 @@
```

```
logging console emergencies
```

```
logging host 192.168.0.1
```

```
!
```

```
-hostname edge01.bjm01
```

```
+hostname edge01.bjm01.lab
```

```
!
```

```
result:
```

```
True
```

Absolute path

Arista device

```
edge01.flw01:
```

```
-----
```

```
already_configured:
```

```
False
```

```
comment:
```

```
diff:
```

```
[edit system]
```

```
- host-name edge01.flw01;
```

```
+ host-name edge01.flw01.lab;
```

```
result:
```

```
True
```

Juniper device

Many vendors, one simple template!

Configuration management

Debug mode

```
$ sudo salt edge01.flw01 net.load_template salt://example.jinja debug=True
```

```
edge01.flw01:
```

```
-----
```

```
already_configured:
```

```
False
```

```
comment:
```

```
diff:
```

```
[edit system]
```

```
- host-name edge01.flw01;
```

```
+ host-name edge01.flw01.lab;
```

```
loaded_config:
```

```
system {
```

```
    host-name edge01.flw01.lab;
```

```
}
```

```
result:
```

```
True
```

Salt path
Template
stored under
the **file_roots**

Debug mode

The result of template rendering.
Not necessarily equal to the diff.


Note: Jinja is painful to debug.
This option is very helpful.
[See more debugging tools](#)

Configuration management


Remote templates

Yes, they can also be elsewhere.
Available options: *salt://*, *ftp://*, *http://*, *https://*,
version control, cloud storage providers etc.

```
$ sudo salt -G 'os:ios' net.load_template http://bit.ly/2gKOj20 peers="['172.17.17.1', '172.17.17.2']"
```



Matches all
devices running
IOS



Loads external template
from <http://bit.ly/2gKOj20>
which shortens the link to
the NAPALM native template for IOS.

Configuration management

Advanced templating: reusing existing data (1)

```
{%- set arp_output = salt.net.arp() -%}  
{%- set arp_table = arp_output['out'] -%}
```

/etc/salt/templates/arp_example.jinja

```
{%- if grains.os|lower == 'iosxr' %} {# if the device is a Cisco IOS-XR #}  
    {%- for arp_entry in arp_table %}  
    arp {{ arp_entry['ip'] }} {{ arp_entry['mac'] }} arpa  
    {%- endfor -%}  
{%- elif grains.vendor|lower == 'juniper' %} {# or if the device is a Juniper #}  
    interfaces {  
        {%- for arp_entry in arp_table %}  
        {{ arp_entry['interface'] }} {  
            family inet {  
                address {{ arp_entry['ip'] }} {  
                    arp {{ arp_entry['ip'] }} mac {{ arp_entry['mac'] }};  
                }  
            }  
        }  
        {%- endfor %}  
    }  
{%- endif %}
```

Retrieving the ARP
table using the
[net.arp](#) function.

Configuration management

Advanced templating: reusing existing data (1)

```
$ sudo salt edge01.flw01 net.load_template salt://arp_example.jinja
```

```
edge01.flw01:
```

```
-----
```

```
already_configured:
```

```
False
```

```
comment:
```

```
diff:
```

```
[edit interfaces xe-0/0/0 unit 0 family inet]
```

```
+ address 10.10.2.2/32 {
```

```
+ arp 10.10.2.2 mac 0c:86:10:f6:7c:a6;
```

```
+ }
```

```
[edit interfaces ae1 unit 1234]
```

```
+ family inet {
```

```
+ address 10.10.1.1/32 {
```

```
+ arp 10.10.1.1 mac 9c:8e:99:15:13:b3;
```

```
+ }
```

```
+ }
```

```
result:
```

```
True
```

Configuration management

Advanced templating: reusing existing data (2)

/etc/salt/templates/route_example.jinja

```
{%- set route_output = salt.route.show('0.0.0.0/0', 'static') -%}
{%- set default_route = route_output['out'] -%}

{%- if not default_route -%} {# if no default route found in the table #}
    {%- if grains.vendor|lower == 'juniper' -%}
routing-options {
    static {
        route 0.0.0.0/0 next-hop {{ pillar.default_route_nh }};
    }
}
    {%- elif grains.os|lower == 'iosxr' -%}
router static address-family ipv4 unicast 0.0.0.0/0 {{ pillar.default_route_nh }}
    {%- endif %}
{%- endif -%}
```

Retrieving the static route data using the route.show function.

This requires appending a new line in the device pillar:

`default_route_nh: 1.2.3.4`

Configuration management

Advanced templating: reusing existing data (2)

```
$ sudo salt 'edge01.oua01' net.load_template salt://route_example.jinja debug=True
```

```
edge01.oua01:
```

```
-----
```

```
already_configured:
```

```
False
```

```
comment:
```

```
diff:
```

```
---
```

```
+++
```

```
@@ -3497,6 +3497,7 @@
```

```
!
```

```
router static
```

```
address-family ipv4 unicast
```

```
+ 0.0.0.0/0 1.2.3.4
```

```
172.17.17.0/24 Null0 tag 100
```

```
loaded_config:
```

```
router static address-family ipv4 unicast 0.0.0.0/0 1.2.3.4
```

```
result:
```

```
True
```

Other simple examples

- Using [postgres.psql_query](#) populate a table in a Postgres database with the network interfaces details (retrieved using [net.interfaces](#))
- Using [bgp.neighbors](#) remove from the BGP config neighbors in *Active* state
- Using [ntp.stats](#), remove unsynchronised NTP peers
- Using [net.environment](#), push high temperature [notifications in Slack](#)

The list can be nearly infinite - depends only on your own use case.

There are thousands of functions already available:

<https://docs.saltstack.com/en/develop/ref/modules/all/index.html>

Note: the examples above are implemented more elegant using states, beacons, reactors, etc.

Real-world orchestration example

State: update NTP peers (1)

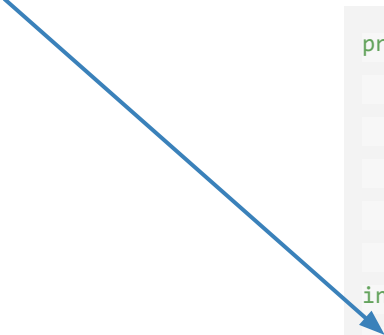
Append directly these lines
in the device pillar, or define
in external file and include:

/etc/salt/pillar/ntp_config.sls

```
ntp.peers:
  - 10.10.1.1
  - 10.10.2.2
ntp.servers:
  - 172.17.17.1
  - 172.17.19.1
```

/etc/salt/pillar/device1.sls

```
proxy:
  proxytype: napalm
  driver: junos
  host: hostname_or_ip_address
  username: my_username
  passwd: my_password
include:
  - ntp_config
```



Better to use the *include*, as
multiple devices can have
the same NTP peers etc.

When including, strip the *.sls*
extension!

Real-world orchestration example

State output: update NTP peers (2)

```
$ sudo salt 'edge01.jnb01' state.sls router.ntp
edge01.jnb01:
-----
ID: update_my_ntp_config
Function: netntp.managed
Result: True
Started: 09:50:41.228728
Duration: 16813.319 ms
Changes:
-----
peers:
-----
removed:
- 10.10.1.1
servers:
-----
added:
- 172.17.17.1
- 172.17.19.1

Summary for edge01.jnb01
-----
Succeeded: 1 (changed=1)
Failed:    0
-----
Total states run:    1
```

Real-world orchestration example

Salt event system

Salt is a [data driven system](#). Each action (job) performed (manually from the CLI or automatically by the system) is uniquely identified and has an identification tag:

```
$ sudo salt-run state.event pretty=True
```

```
salt/job/20170110130619367337/new {
```

```
  "_stamp": "2017-01-10T13:06:19.367929",
```

```
  "arg": [],
```

```
  "fun": "probes.results",
```

```
  "jid": "20170110130619367337",
```

```
  "minions": [
```

```
    "edge01.bjm01"
```

```
  ],
```

```
  "tgt": "edge01.bjm01",
```

```
  "tgt_type": "glob",
```

```
  "user": "mircea"
```

```
}
```

Unique job tag



Real-world orchestration example

Reactor

Using the job tags, you can identify events (triggers) and react (action):

/etc/salt/master

```
reactor:  
  - 'salt/job/*/ret/*':  
    - salt://reactor_example.sls
```

Unique job tags (regular expression): in this example will match any job returns

When this event occurs, execute this reactor descriptor.

/etc/salt/reactors/reactor_example.sls

```
invoke_orchestrate_file:  
  runner.state.orchestrate:  
    - mods: orch.do_complex_thing  
    - pillar:  
      event_tag: {{ tag }}  
      event_data: {{ data | json() }}
```

Real-world orchestration example


Beacon example: inotify

Beacons let you use the Salt event system to monitor non-Salt processes.

/etc/salt/proxy

```
beacons:  
  inotify:  
    /etc/salt/pillar/ntp_config.sls:  
      mask:  
        - modify  
      disable_during_state_run: True
```

Will fire an event when
updating
/etc/salt/pillar/ntp_config.sls



Real-world orchestration example

Beacon event tag example

This event is fired when a change is made and saved to `/etc/salt/pillar/ntp_config.sls`:

```
salt/beacon/device1/inotify//etc/salt/pillar/ntp_config.sls {
  "_stamp": "2017-01-09T15:59:37.972753",
  "data": {
    "change": "IN_IGNORED",
    "id": "device1",
    "path": "/etc/salt/pillar/ntp_config.sls"
  },
  "tag": "salt/beacon/device1/inotify//etc/salt/pillar/ntp_config.sls"
}
```

Using the reactor system, one can match these event tags and take actions when they happen.

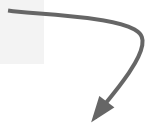
Real-world orchestration example

Beacon event tag example

React when the `/etc/salt/pillar/ntp_config.sls` is changed

`/etc/salt/master`

```
reactor:  
  - 'salt/beacon/*/inotify//etc/salt/pillar/ntp_config.sls':  
    - salt://run_ntp_state_on_pillar_update.sls
```



`/etc/salt/reactors/run_ntp_state_on_pillar_update.sls`

```
run_ntp_state:  
  local.state.sls:  
    - tgt: {{ data['id'] }}  
    - arg:  
      - router.ntp
```

This is how the reactor system knows that a state execution is required.

Run the state against the minion ID that triggered the event

Run the ntp state defined earlier (slides #35-#36).

Real-world orchestration example

Beacon event tag example

... and that's it!


From now on, whenever you update */etc/salt/pillar/ntp_config.sls*,
it will automatically update your routers' config.

And you maintain entities of data, not pseudo-formatted text files,
regardless on the device vendor.

Advanced topics

Mine

Embedded caching



Read more: <https://docs.saltstack.com/en/latest/topics/mine/>

Advanced topics

The Salt API

You can also execute commands remotely, via HTTPS

Easy to setup, easy to use

/etc/salt/master

```
rest_cherry.py:  
  port: 8001  
  ssl_crt: /etc/nginx/ssl/my_certificate.pem  
  ssl_key: /etc/nginx/ssl/my_key.key
```



```
curl -sSk  
https://salt-master-ns-or-ip:8001/run \  
-H 'Content-type: application/json' \  
-d ' [{  
    "client": "local",  
    "tgt": "<target>",  
    "fun": "net.arps",  
    "username": "my username",  
    "password": "my password",  
    "eauth": "pam"  
}] '
```

Sources

<https://github.com/mirceaulinic/talks/tree/master/APRICOT2017/sources>

More advanced topics

- Orchestration: define complex workflows
<https://docs.saltstack.com/en/latest/topics/orchestrate/index.html>
- Publish events to external services (e.g.: logstash, hipchat)
<https://docs.saltstack.com/en/develop/ref/engines/all/index.html>
- Pillar: load data from external services, not just static
<https://docs.saltstack.com/en/develop/ref/pillar/all/>
- Custom authentication methods for the minions
<https://docs.saltstack.com/en/develop/ref/auth/all/index.html>
- Forward outputs in external data systems on runtime
<https://docs.saltstack.com/en/develop/ref/returners/all/index.html>

How can you contribute?

- NAPALM Automation:
<https://github.com/napalm-automation>
- SaltStack
<https://github.com/saltstack/salt>

Need help/advice?

Join [#saltstack #napalm](https://networktoencode.herokuapp.com/rooms)

By email:

- Mircea Ulinic: mircea@cloudflare.com
- Jerome Fleury: jf@cloudflare.com

Questions



References

[Authentication system](#)

[Beacons](#)

[Engines](#)

[Event System](#)

[Grains](#)

[Jinja](#)

[load_template documentation](#)

[Master config file, example](#)

[Master configuration options](#)

[Mine](#)

[NAPALM](#)

[NAPALM BGP execution module functions](#)

[NAPALM Grains](#)

[NAPALM Installation](#)

[NAPALM network execution module functions](#)

[NAPALM NTP execution module functions](#)

[NAPALM Proxy](#)

[NAPALM route execution module functions](#)

[Nested outputter](#)

[NETAPI Modules](#)

[Netconfig state](#)

[Node Groups](#)

[NTP state](#)

[Orchestration](#)

[Pillar](#)

[Pillar modules](#)

[Proxy Minion](#)

[Reactor](#)

[REST CherryPy](#)

[Returners](#)

[Runners](#)

[Salt 2016.11 \(Carbon\) release notes](#)

[Salt Get Started](#)

[Salt Installation](#)

[Salt Walkthrough](#)

[SaltStack Package Repo](#)

[SNMP state](#)

[States](#)

[Targeting minions](#)

[The Top file](#)

[Users state](#)

[YAML](#)